

OpenMP mini-app practical

Ben Cumming

Swiss National Supercomputing Center (CSCS)



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

CSCS

Swiss National Supercomputing Centre



Adding OpenMP to the mini-app

- The aim of this practical is to take the serial mini-app implementation, and make it faster with OpenMP
 - while still getting the correct answer

Hints

- Before starting find two sets of input parameters that converge for the serial version
 - note the time to solution
 - you want this to get faster as you add OpenMP
 - **NOTE:** but you might have to add quite a few directives before things actually get faster
 - note the number of conjugate gradient iterations
 - use this to check after you add each directive that you are still getting the right answer
 - **NOTE:** remember that there will be some small variations because floating point operations are not commutative



First test

- Get the code, by checking it out from github

```
> git pull
> cd <SummerSchool2015path>/miniapp/openmp
> ls
cxx fortran
> cd cxx
> make
> aprun ./main 128 128 100 0.01
<note time to solution and conjugate gradient iterations>
> aprun ./main 256 256 200 0.01
<note time to solution and conjugate gradient iterations>
```

I choose the C++ version here

Step 1

- replace the welcome message in main.cpp/main.f90 with a message that tells the user that
 - this is the openmp version
 - how many threads it is using

```
> make
> OMP_NUM_THREADS=8 aprun ./main -cc none 256 256 200 0.01
...
=====
Welcome to mini-stencil!
version :: Fortran90 OpenMP with 8 threads
...
```

Step 2: Linear Algebra

- Open [linalg.cpp/f90](#) and add directives to the functions subroutines `ss_XXXX`
 - do one or two at a time
 - recompile frequently and run with 8 threads to check that you are still getting the right answer
- Once finished with that file, did your changes make any improvement?
 - compare the 128x128 and 256x256 results

Step 3: the diffusion stencil

- The final step is to parallelize the stencil operator in `operators.cpp/f90`.
- The nested for/do loop is an obvious target
 - it covers $n_x \times n_y$ grid points
- How about the boundary loops?



Step 4: testing

- how does it scale at different resolutions?
 - 32x32
 - 64x64
 - 128x128
 - 256x256
 - 512x512
 - 1024x1024
- Advanced C++:
 - can you implement first touch memory allocation in the C++ version?



Extras

- can you implement first touch memory allocation
 - requires adding just one OpenMP directive in C++
- does the stencil kernel vectorize?
 - look at Cray and Intel compiler vectorization reports
 - can you make it vectorize?



Thanks for your attention



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich


CSCS

Swiss National Supercomputing Centre



Code Walkthrough

- There are three modules of interest
 - [main.f90/main.cpp](#) : initialization and main time stepping loop
 - [linalg.f90/linalg.cpp](#) : the BLAS level 1 (vector-vector) kernels and conjugate gradient solver
 - [operators.f90/operators.cpp](#) : the stencil operator for the finite volume discretization



the vector-vector kernels and diffusion operator are the only kernels that have to be parallelized

parameter notes

- Compile

```
> make
```

- Run interactively (use salloc beforehand)

```
> aprun ./main 128 128 100 0.01
```

- the grid is 128 x 128 grid points
- take 100 time steps
- run simulation for t=0.01

- Or run batch job

```
> sbatch job.daint  
... when job is finished ...  
> cat job.out
```



Output

```
=====
version :
mesh
time :
=====

step 1 required 4 iterations for residual 7.21951e-07
step 2 required 4 iterations for residual 7.9975e-07
...
step 99 required 12 iterations for residual 7.21951e-07
step 100 required 12 iterations for residual 7.21951e-07
-----
simulation took 1.58408 seconds
8127 conjugate gradient iterations, at rate of 5130.43 iters/second
920 newton iterations
-----
Goodbye!
```

The number of conjugate gradient iterations, which should always be constant for a given mesh size and time parameters. Can be used to check that changes to the code are still getting the correct result. There will be small variations due to the imprecise nature of floating point operations.

time to solution

best way to compare different implementations



Visualize the answer

- The application generates two data files with the final solution: `output.bin` and `output.bov`
- There is a script to draw a contour plot of the solution

```
> ls output.*  
output.bin output.bov  
> module load python/2.7.6  
> python plotting.py
```

requires X-windowing
make sure you connect with "ssh -X"

