



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Summer School 2015

MPI Profiling

<http://github.com/eth-cscs/SummerSchool2015/wiki> - DAY4

July 2015

# Summary

- Parallel performance
- Profiling the MPI mini app
  - with Cray's perftools-lite
  - with scorep



**CSCS**


Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# Parallel Performance

---

## Scalability results (C++/MPI miniapp)



## Getting timings

```
course51@daint103:~ /usr/bin/time -p aprun -n 32 256 256 400 0.1
=====
Welcome to mini-stencil!
version      :: with MPI : 32 MPI ranks
mesh        :: 256 * 256 dx = 0.00392157
time        :: 400 time steps from 0 .. 0.1
...
----- Goodbye!
real 5.41
```

<===== real time in seconds



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# WHY ?

---



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# Perftools

---

# MPI miniapp + perftools: compiling

## Compiling with the tool

- Cray's perftools supports all 4 compilers
- module load `perftools-lite`; man `perftools-lite`
- make clean; make

## Recompile with optimisation flags on

```
course51@daint103:~ make
CC -O3 -c stats.cpp -o stats.o
CC -O3 -c data.cpp -o data.o
CC -O3 -c operators.cpp -o operators.o
CC -O3 -c linalg.cpp -o linalg.o
CC -O3 *.o main.cpp -o main.exe

INFO: creating the CrayPat-instrumented exec 'main.exe'
      (sample_profile) ...OK

INFO: A maximum of 53 functions from group 'io' will be traced.
INFO: A maximum of 292 functions from group 'mpi' will be traced.
INFO: A maximum of 22 func. from group 'realtime' will be traced.
INFO: A maximum of 54 func. from group 'syscall' will be traced.
```

# MPI miniapp + perftools: running

Run to get the performance report (4 CN here)

```
course51@daint103:~ aprun -n 32 256 256 400 0.1
```

```
=====
Welcome to mini-stencil!
```

```
...
```

```
----- Goodbye!
```


```
Table 1: Profile by Function Group and Function
         (top 10 functions shown)
```

| Samp%  | Samp  | Imb. | Imb.  | Group           |
|--------|-------|------|-------|-----------------|
|        |       | Samp | Samp% | Function        |
|        |       |      |       | PE=HIDE         |
| 100.0% | 444.3 | --   | --    | Total           |
| 47.3%  | 210.2 | --   | --    | MPI             |
| 34.2%  | 151.9 | 29.1 | 16.6% | MPI_Allreduce   |
| 6.1%   | 27.1  | 17.9 | 41.1% | MPI_Cart_create |
| 3.2%   | 14.4  | 14.6 | 51.8% | MPI_Isend       |
| 3.0%   | 13.4  | 10.6 | 45.4% | MPI_Waitall     |
| 40.2%  | 178.8 | --   | --    | USER            |
| 7.7%   | 34.1  | --   | --    | OMP             |
| 4.8%   | 21.2  | --   | --    | ETC             |

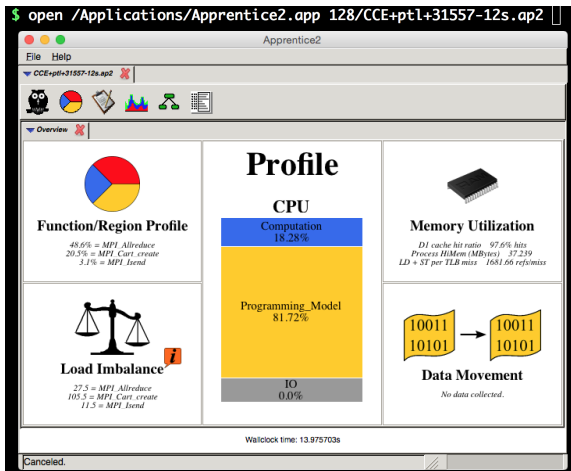


# MPI miniapp + perftools: Because !

## Scalability results (C++/MPI miniapp): 256x256



# MPI miniapp + perftools: apprentice2




- You can install app2 on your laptop: `/opt/cray/perftools/default/share/desktop_installers/`

# MPI miniapp + perftools: tracing

## app2

- module load `perftools-lite`; man `perftools-lite`
- export `CRAYPAT_LITE=event_profile`
- make clean; make; aprun ...





**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# Scorep

---

# MPI miniapp + scorep: compiling

## Compiling with the tool

- module rm perftools-lite
- module load scorep; module help scorep
- make clean; make CXX='scorep --mpp=mpi CC'

## Recompile with optimisation flags on

```
course51@daint103:~ make CXX="scorep --mpp=mpi CC" CXXFLAGS="-O3 -h
noomp"
scorep --mpp=mpi CC -O3 -h noomp -c stats.cpp
scorep --mpp=mpi CC -O3 -h noomp -c data.cpp
scorep --mpp=mpi CC -O3 -h noomp -c operators.cpp
scorep --mpp=mpi CC -O3 -h noomp -c linalg.cpp
scorep --mpp=mpi CC -O3 -h noomp *.o main.cpp -o main
```

# MPI miniapp + scorep: running

## Run to get the performance report dir

```
course51@daint103:~ aprun -n 8 CCE+sc141 128 128 25 0.0001
=====
Welcome to mini-stencil!
...
----- Goodbye!

New dir ==> scorep-20150722_1150_1904915555376033/
```

## View the text performance report

```
course51@daint103:~ scorep-score \
scorep-20150722_1150_1904915555376033/profile.cubex
```

|  | type   | max_buf[B]  | visits      | time[s] | time[%] | time/visit[us] |     |
|--|--------|-------------|-------------|---------|---------|----------------|-----|
|  | region |             |             |         |         |                |     |
|  | ALL    | 501,600,399 | 166,898,392 | 84.55   | 100.0   | 0.51           | ALL |
|  | USR    | 501,322,776 | 166,859,984 | 60.97   | 72.1    | 0.37           | USR |
|  | MPI    | 247,839     | 28,480      | 0.72    | 0.9     | 25.31          | MPI |
|  | COM    | 29,784      | 9,928       | 22.85   | 27.0    | 2301.87        | COM |

# MPI miniapp + scorep: scalasca


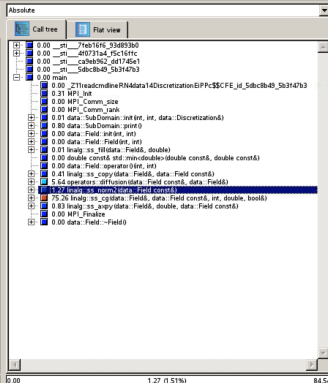
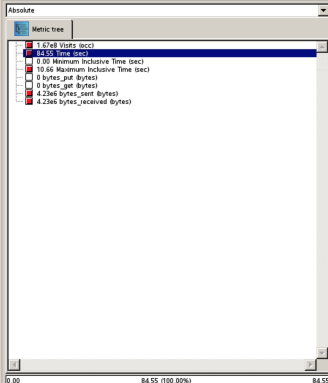
xterm

```
piccinal@santis01:/scratch/santis/piccinal/summer-school.git/mpi/cxx $ square scorep-20150722_1150_1904915555376033/profile.  
INFO: Displaying scorep-20150722_1150_1904915555376033/profile.cubex...
```

Cube-4.3.1: scorep-20150722\_1150\_1904915555376033/profile.cubex

File Display Help

Restore Setting Save Settings




Selected 'linalg: ss\_norm2 (data: Field const&)'

# MPI miniapp + scorep: vampir (1)


## vampir

- module rm perftools-lite
- module load scorep; module help scorep
- make clean; make;
- export SCOREP\_ENABLE\_TRACING=true; aprun ...






# MPI miniapp + perftools: vampir (2)



## MPI miniapp + perftools: vampir (3)





**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# Reveal

---

# serial miniapp + reveal

## Compiling with the tool (CCE compiler only)

- Cray reveal can help you add OpenMP directives in the src:
- module load perftools
- make clean; make FFLAGS="-O3 -eZ `-hpl=reveal.pl`"
- no need to run
- reveal reveal.pl # requires X

# serial miniapp + reveal: loops

|  |   |
|--|---|
| <pre> 302 do iter=1, maxiters 303     Ap = A*p 304   call ss_lcomb(v, one, xold, eps, p, N) 305   call diffusion(v, fx) 306   call ss_scaled_diff(Ap, eps_inv, fx, Fxold, N) 307 308     alpha = rold / p*Ap 309   alpha = rold / ss_dot(p, Ap, N) 310 311     x = alpha*p 312   call ss_axpy(x, alpha, p, N) 313 314     r = -alpha*Ap 315   call ss_axpy(r, -alpha, Ap, N) 316 317     find new norm 318   rnew = ss_dot(r, r, N) 319 320     test for convergence 321   if( dabs(rnew-xtol) ) then 322     success = .true. 323     exit 324   endif 325 326     p = r + rnew*rold * p 327   call ss_lcomb(p, one, r, rnew/rold, p, N) 328 329   rold = rnew </pre>                           | <pre> 35   the interior grid points 36 do j = 2, jend 37   do i = 2, iend 38     s(i,j) = -(4.*alpha) * u(i,j) 39             + u(i-1,j) + u(i+1,j) 40             + u(i,j-1) + u(i,j+1) 41             + alpha*x_oldd(i,j) 42             + dxs*u(i,j)*(1.0-0.8 - u(i,j)) 43 44   end do 45 end do 46 47   the east boundary 48 i = options\$Nnx 49 do j = 2, jend 50   s(i,j) = -(4.*alpha) * u(i,j) 51             + u(i-1,j) + u(i,j-1) + u(i,j+1) 52             + alpha*x_oldd(i,j) + bndE(j) 53             + dxs*u(i,j)*(1.0-0.8 - u(i,j)) 54 55 end do 56 57   the west boundary 58 i = 1 59 do j = 2, jend 60   s(i,j) = -(4.*alpha) * u(i,j) 61             + u(i+1,j) + u(i,j-1) + u(i,j+1) 62             + alpha*x_oldd(i,j) + bndW(j) 63             + dxs*u(i,j)*(1.0-0.8 - u(i,j)) 64 65 end do </pre>   |
| <pre> linalg.f90 383,5 121 do timestep = 1, nt 122     set x_new and x_old to be the solution 123   call ss_copy(x_old, x_new, N) 124 125   converged = .false. 126   do it = 1, 50 127       compute residual : requires both x_new and x_old 128     call diffusion(x_new, b) 129     residual = ss_norm2(b, N) 130 131       check for convergence 132     if(residual&lt;tolerance) then 133       converged = .true. 134       exit 135     endif 136 137       solve linear system to get -delta x 138     call ss_cg\$delta, b, 200, tolerance, cg_converged 139 140       check that the CG solver converged 141     if(.NOT. cg_converged) then 142       exit 143     endif 144 </pre> | <pre> 96linoperators.f90 37,5   Loop   Loop   Time   Loop   Loop   Loop   Loop   Loop   Function~/LOOP[.]   TimeK   Time   Adj.   Hit   Trips   Trips   Trips     Avg   Min   Max   ----- 11 99.4K   0.847536   0.000056   1   200.0   200   200   idiffusion_serial_LOOP.3.11.121 12 98.6K   0.848804   0.000661   200   50.0   50   50   idiffusion_serial_LOOP.4.11.126 21 70.6K   0.681916   0.004540   405   4.4   2   13   iss_cg\$linalg_LOOP.1.11.382 3a1 30.8K   0.263112   0.008732   3187   254.0   254   254   idiffusionoperators_LOOP.1.11.36 3a1 29.8K   0.254380   0.254380   809498   254.0   254   254   idiffusionoperators_LOOP.2.11.37 3b1 16.4K   0.139561   0.139561   3949   65536.0   65536   65536   iss_cg\$linalg_LOOP.1.11.115 16.0K   0.136134   0.136134   3139   65536.0   65536   65536   ss_lcomb\$linalg_LOOP.1.11.206 17.0K   0.092907   0.092907   3949   65536.0   65536   65536   ss_dot\$linalg_LOOP.1.11.47 18.4K   0.088793   0.088793   1772   65536.0   65536   65536   ss_scaled_diff\$linalg_LOOP.1.11.160 19.0K   0.032996   0.032996   1010   65536.0   65536   65536   ss_copy\$linalg_LOOP.1.11.226 20.1K   0.026339   0.026339   405   65536.0   65536   65536   ss_add_scaled_diff\$linalg_LOOP.1.11.138 21.0K   0.015104   0.015104   810   65536.0   65536   65536   ss_fill\$linalg_LOOP.1.11.93 22.0K   0.014589   0.014589   405   65536.0   65536   65536   ss_scale\$linalg_LOOP.1.11.182 23.0K   0.009565   0.009565   605   65536.0   65536   65536   ss_norm2\$linalg_LOOP.1.11.70 </pre> |

# serial miniapp + reveal: scoping

The screenshot displays the Reveal IDE interface. The main window shows a Fortran source file with the following code:

```
113 ! the logic
114 do i = 1, N
115   y(i) = alpha*x(i) + y(i)
116 enddo
117
118 ! update the flops counter
119 flops_blas1 = flops_blas1 + 2*N
120 end subroutine
121
122 ! computes y = x + alpha*(l..r)
123 ! y, x, l and r are vectors of length N
124 ! alpha is a scalar
125 subroutine ss_add_scaled_diff(y, x, alpha, l, r, N)
126 ! arguments
127 real (kind=8), intent(in) :: alpha
128 real (kind=8), intent(in) :: x(N)
129 real (kind=8), intent(inout) :: y(N)
```

The left sidebar shows a navigation tree with the following structure:

- 1.05% DIFFUSION\_SERIAL
  - DIFFUSION
    - CG\_INIT
    - SS\_DOT
    - SS\_NORM2
    - SS\_FILL
    - SS\_AXPY
    - Loop@115
    - SS\_ADD\_SCALED\_DIFF
    - SS\_SCALED\_DIFF
    - SS\_SCALE
    - SS\_LCOMB
    - SS\_COPY
    - SS\_CG
    - SWAP\_DATA
    - READCDLINE
    - ERROR

The right sidebar shows the "Reveal OpenMP Scoping" window with the following table:

| Reveal OpenMP Scoping |        |                 |      |
|-----------------------|--------|-----------------|------|
| Scope Loops           |        | Scoping Results |      |
| linalg.f90: Loop@115  |        |                 |      |
| Name                  | Type   | Scope           | Info |
| i                     | Scalar | Private         |      |
| alpha                 | Scalar | Shared          |      |
| n                     | Scalar | Shared          |      |
| x                     | Array  | Shared          |      |
| y                     | Array  | Shared          |      |

The bottom status bar shows the following information:

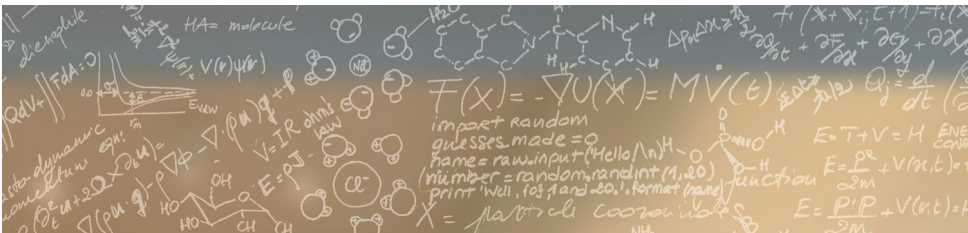
- Info - Line 115
- A loop starting at line 115 was unrolled 2 times.
- A loop starting at line 115 was vectorized.



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



Thank you for your attention.